

# Neural Network Library Documentation

## Overview

This documentation provides an overview of a neural network library, covering the main components required to build, train, and evaluate a neural network. The library includes classes for data loading, network layers, activation functions, dropout, loss functions, and optimizers.

## Network

### Class: Network

The `Network` class represents the neural network, managing layers, forward and backward passes, and the overall training and evaluation process.

- **Constructor: `__init__(self, loss, optimizer)`**
  - **Parameters:**
    - \* `loss` (class): Loss function class (e.g., `MSELoss`, `MAELoss`, `BCELoss`).
    - \* `optimizer` (object): Optimizer instance (e.g., `SGD`, `Adadelata`, `Adam`).
- **Method: `train_mode(self)`**
  - **Description:** Sets the network to training mode.
- **Method: `evaluation_mode(self)`**
  - **Description:** Sets the network to evaluation mode.
- **Method: `add_layer(self, layer, layer_inputs, layer_outputs, distribution="normal")`**
  - **Description:** Adds a layer to the network.
  - **Parameters:**
    - \* `layer` (class): Layer class (e.g., `Linear`).
    - \* `layer_inputs` (int): Number of input neurons to the layer.
    - \* `layer_outputs` (int): Number of output neurons from the layer.
    - \* `distribution` (str): Distribution for weight initialization. Default is "normal".
- **Method: `add_dropout(self, layer_inputs, fraction=0.5)`**
  - **Description:** Adds a dropout layer to the network.
  - **Parameters:**
    - \* `layer_inputs` (int): Number of input neurons to the layer.
    - \* `fraction` (float): Fraction of the input units to drop. Default is 0.5.
- **Method: `add_activation(self, activation)`**
  - **Description:** Adds an activation function to the network.
  - **Parameters:**
    - \* `activation` (object): Activation function instance (e.g., `Sigmoid`, `Tanh`, `Relu`, `LeakyRelu`).
- **Method: `evaluate(self, x)`**
  - **Description:** Performs a forward pass through the network for evaluation.
  - **Parameters:**
    - \* `x` (numpy array): Input data.
  - **Returns:** Output of the network.
- **Method: `train(self, x, y)`**
  - **Description:** Performs a forward and backward pass through the network for training.
  - **Parameters:**
    - \* `x` (numpy array): Input data.
    - \* `y` (numpy array): Correct labels.
  - **Returns:** Mean loss of the batch.

### Example Usage

```
import numpy as np
from Network import Network
from Linear import Linear
from ActivationFunctions import Relu, Sigmoid
from Dropout import Dropout
from LossFunctions import MSELoss
from Optimizers import Adam

# Example input and labels
```

```

x = np.random.rand(10, 5)
y = np.random.rand(10, 1)

# Initialize optimizer
optimizer = Adam(eta=0.001)

# Initialize network
network = Network(loss=MSELoss, optimizer=optimizer)

# Add layers
network.add_layer(Linear, layer_inputs=5, layer_outputs=10, distribution="normal_xavier")
network.add_activation(Relu())
network.add_dropout(layer_inputs=10, fraction=0.5)
network.add_layer(Linear, layer_inputs=10, layer_outputs=1, distribution="normal_xavier")
network.add_activation(Sigmoid())

# Set to training mode
network.train_mode()

# Train the network
loss = network.train(x, y)
print("Training Loss:", loss)

# Set to evaluation mode
network.evaluation_mode()

# Evaluate the network
output = network.evaluate(x)
print("Evaluation Output:", output)

```

## Data Loader

### Class: DataLoader

The DataLoader class handles loading and iterating over datasets for training and evaluation.

- **Constructor: `__init__(self, inputs, labels, batch_size=1, shuffle=False, drop_last=False)`**
  - **Parameters:**
    - \* `inputs` (numpy array): Input data.
    - \* `labels` (numpy array): Corresponding labels.
    - \* `batch_size` (int): Size of each data batch. Default is 1.
    - \* `shuffle` (bool): Whether to shuffle the data before each epoch. Default is `False`.
    - \* `drop_last` (bool): Whether to drop the last batch if it's smaller than the batch size. Default is `False`.
- **Method: `__iter__(self)`**
  - **Description:** Creates an iterator over the data.
  - **Returns:** Batches of inputs and labels.

### Example Usage

```

import numpy as np
from DataLoader import DataLoader

# Example input and labels
inputs = np.random.rand(100, 20)
labels = np.random.rand(100, 1)

# Initialize data loader
data_loader = DataLoader(inputs, labels, batch_size=10, shuffle=True)

# Iterate over batches
for batch_inputs, batch_labels in data_loader:

```

```
print("Batch inputs:", batch_inputs)
print("Batch labels:", batch_labels)
```

## Layers

### Class: Linear

The Linear class defines a fully connected layer in the neural network.

- **Constructor: `__init__(self, input_neurons, output_neurons, optimizer, distribution)`**
  - **Parameters:**
    - \* `input_neurons` (int): Number of input neurons.
    - \* `output_neurons` (int): Number of output neurons.
    - \* `optimizer` (object): Optimizer instance (e.g., SGD, Adadelta, Adam).
    - \* `distribution` (str): Distribution for weight initialization (e.g., "normal\_xavier", "uniform\_xavier", "normal").
- **Method: `forward(self, x)`**
  - **Description:** Performs the forward pass.
  - **Parameters:**
    - \* `x` (numpy array): Input data.
  - **Returns:** Output of the layer.
- **Method: `backward(self, gradient, layer_input)`**
  - **Description:** Performs the backward pass.
  - **Parameters:**
    - \* `gradient` (numpy array): Gradient from the next layer.
    - \* `layer_input` (numpy array): Input data to the layer during forward pass.
  - **Returns:** Gradient for the previous layer.

### Example Usage

```
import numpy as np
from Linear import Linear
from Optimizers import Adam

# Example input
x = np.random.rand(10, 5)

# Initialize optimizer
optimizer = Adam(eta=0.001)

# Initialize linear layer
linear_layer = Linear(input_neurons=5, output_neurons=3, optimizer=optimizer, distribution="normal_xavier")

# Forward pass
output = linear_layer.forward(x)
print("Layer output:", output)

# Backward pass
gradient = np.random.rand(10, 3)
backprop_gradient = linear_layer.backward(gradient, x)
print("Backpropagation gradient:", backprop_gradient)
```

## Activation Functions

### Class: Sigmoid

The Sigmoid class implements the sigmoid activation function.

- **Method: `forward(self, layer_input)`**
  - **Description:** Applies the sigmoid function.
  - **Parameters:**
    - \* `layer_input` (numpy array): Input data.

- **Returns:** Activated output.
- **Method: backward(self, gradient, layer\_input)**
  - **Description:** Computes the gradient of the sigmoid function.
  - **Parameters:**
    - \* **gradient** (numpy array): Gradient from the next layer.
    - \* **layer\_input** (numpy array): Input data to the layer during forward pass.
  - **Returns:** Gradient for the previous layer.

#### Class: Tanh

The Tanh class implements the tanh activation function.

- **Method: forward(self, layer\_input)**
  - **Description:** Applies the tanh function.
  - **Parameters:**
    - \* **layer\_input** (numpy array): Input data.
  - **Returns:** Activated output.
- **Method: backward(self, gradient, layer\_input)**
  - **Description:** Computes the gradient of the tanh function.
  - **Parameters:**
    - \* **gradient** (numpy array): Gradient from the next layer.
    - \* **layer\_input** (numpy array): Input data to the layer during forward pass.
  - **Returns:** Gradient for the previous layer.

#### Class: Relu

The Relu class implements the ReLU activation function.

- **Method: forward(self, layer\_input)**
  - **Description:** Applies the ReLU function.
  - **Parameters:**
    - \* **layer\_input** (numpy array): Input data.
  - **Returns:** Activated output.
- **Method: backward(self, gradient, layer\_input)**
  - **Description:** Computes the gradient of the ReLU function.
  - **Parameters:**
    - \* **gradient** (numpy array): Gradient from the next layer.
    - \* **layer\_input** (numpy array): Input data to the layer during forward pass.
  - **Returns:** Gradient for the previous layer

#### Class: LeakyRelu

The LeakyRelu class implements the leaky ReLU activation function.

- **Constructor: \_\_init\_\_(self, a=0.01)**
  - **Parameters:**
    - \* **a** (float): Slope of the activation function for negative inputs. Default is 0.01.
- **Method: forward(self, layer\_input)**
  - **Description:** Applies the leaky ReLU function.
  - **Parameters:**
    - \* **layer\_input** (numpy array): Input data.
  - **Returns:** Activated output.
- **Method: backward(self, gradient, layer\_input)**
  - **Description:** Computes the gradient of the leaky ReLU function.
  - **Parameters:**
    - \* **gradient** (numpy array): Gradient from the next layer.
    - \* **layer\_input** (numpy array): Input data to the layer during forward pass.
  - **Returns:** Gradient for the previous layer.

## Example Usage

```
import numpy as np
from ActivationFunctions import Sigmoid, Tanh, Relu, LeakyRelu

# Example input
x = np.random.rand(10, 5)

# Sigmoid activation
sigmoid = Sigmoid()
sigmoid_output = sigmoid.forward(x)
print("Sigmoid output:", sigmoid_output)

# Tanh activation
tanh = Tanh()
tanh_output = tanh.forward(x)
print("Tanh output:", tanh_output)

# ReLU activation
relu = Relu()
relu_output = relu.forward(x)
print("ReLU output:", relu_output)

# Leaky ReLU activation
leaky_relu = LeakyRelu(a=0.1)
leaky_relu_output = leaky_relu.forward(x)
print("Leaky ReLU output:", leaky_relu_output)
```

## Dropout

Class: Dropout

The Dropout class implements dropout regularization.

- **Constructor: `__init__(self, input_neurons, fraction, training_mode)`**
  - **Parameters:**
    - \* `input_neurons` (int): Number of input neurons.
    - \* `fraction` (float): Fraction of the input units to drop.
    - \* `training_mode` (list): List containing a single boolean indicating training mode.
- **Method: `forward(self, layer_input)`**
  - **Description:** Applies dropout during training.
  - **Parameters:**
    - \* `layer_input` (numpy array): Input data.
  - **Returns:** Output after applying dropout.
- **Method: `backward(self, gradient, *)`**
  - **Description:** Computes the gradient for dropout.
  - **Parameters:**
    - \* `gradient` (numpy array): Gradient from the next layer.
  - **Returns:** Gradient for the previous layer.

## Example Usage

```
import numpy as np
from Dropout import Dropout

# Example input
x = np.random.rand(10, 5)

# Initialize dropout
training_mode = [True]
dropout = Dropout(input_neurons=5, fraction=0.5, training=training_mode)
```

```

# Forward pass (training mode)
dropout_output = dropout.forward(x)
print("Dropout output (training):", dropout_output)

# Set to evaluation mode
dropout.training[0] = False
dropout_output_eval = dropout.forward(x)
print("Dropout output (evaluation):", dropout_output_eval)

```

## Loss Functions

### Class: MSELoss

The MSELoss class implements the Mean Squared Error loss function.

- **Method: `loss(self, y_pred, y_true)`**
  - **Description:** Computes the MSE loss.
  - **Parameters:**
    - \* `y_pred` (numpy array): Predicted values.
    - \* `y_true` (numpy array): True values.
  - **Returns:** MSE loss.
- **Method: `loss_gradient(self, y_pred, y_true)`**
  - **Description:** Computes the gradient of the MSE loss.
  - **Parameters:**
    - \* `y_pred` (numpy array): Predicted values.
    - \* `y_true` (numpy array): True values.
  - **Returns:** Gradient of the MSE loss.

### Class: MAELoss

The MAELoss class implements the Mean Absolute Error loss function.

- **Method: `loss(self, y_pred, y_true)`**
  - **Description:** Computes the MAE loss.
  - **Parameters:**
    - \* `y_pred` (numpy array): Predicted values.
    - \* `y_true` (numpy array): True values.
  - **Returns:** MAE loss.
- **Method: `loss_gradient(self, y_pred, y_true)`**
  - **Description:** Computes the gradient of the MAE loss.
  - **Parameters:**
    - \* `y_pred` (numpy array): Predicted values.
    - \* `y_true` (numpy array): True values.
  - **Returns:** Gradient of the MAE loss.

### Class: BCELoss

The BCELoss class implements the Binary Cross-Entropy loss function.

- **Method: `loss(self, y_pred, y_true)`**
  - **Description:** Computes the BCE loss.
  - **Parameters:**
    - \* `y_pred` (numpy array): Predicted values.
    - \* `y_true` (numpy array): True values.
  - **Returns:** BCE loss.
- **Method: `loss_gradient(self, y_pred, y_true)`**
  - **Description:** Computes the gradient of the BCE loss.
  - **Parameters:**
    - \* `y_pred` (numpy array): Predicted values.
    - \* `y_true` (numpy array): True values.
  - **Returns:** Gradient of the BCE loss.

## Example Usage

```
import numpy as np
from LossFunctions import MSELoss, MAELoss, BCELoss

# Example predictions and true values
y_pred = np.random.rand(10, 1)
y_true = np.random.rand(10, 1)

# MSE Loss
mse_loss = MSELoss()
loss_value = mse_loss.loss(y_pred, y_true)
loss_gradient = mse_loss.loss_gradient(y_pred, y_true)
print("MSE Loss value:", loss_value)
print("MSE Loss gradient:", loss_gradient)

# MAE Loss
mae_loss = MAELoss()
loss_value = mae_loss.loss(y_pred, y_true)
loss_gradient = mae_loss.loss_gradient(y_pred, y_true)
print("MAE Loss value:", loss_value)
print("MAE Loss gradient:", loss_gradient)

# BCE Loss
bce_loss = BCELoss()
loss_value = bce_loss.loss(y_pred, y_true)
loss_gradient = bce_loss.loss_gradient(y_pred, y_true)
print("BCE Loss value:", loss_value)
print("BCE Loss gradient:", loss_gradient)
```

## Optimizers

### Class: SGD

The SGD class implements the Stochastic Gradient Descent optimizer.

- **Constructor: `__init__(self, learning_rate=0.001, momentum=0, nesterov=False)`**
  - **Parameters:**
    - \* `learning_rate` (float): Learning rate. Default is 0.001.
    - \* `momentum` (float): Momentum factor. Default is 0.
    - \* `nesterov` (bool): Whether to use Nesterov momentum. Default is False.
- **Method: `optimize(self, weights, biases, gradient_weights, gradient_biases)`**
  - **Description:** Updates weights and biases based on gradients.
  - **Parameters:**
    - \* `weights` (numpy array): Current weights.
    - \* `biases` (numpy array): Current biases.
    - \* `gradient_weights` (numpy array): Gradients of weights.
    - \* `gradient_biases` (numpy array): Gradients of biases.
  - **Returns:** Updated weights and biases.

### Class: Adadelata

The Adadelata class implements the Adadelata optimizer.

- **Constructor: `__init__(self, gamma=0.9, eps=1e-8)`**
  - **Parameters:**
    - \* `gamma` (float): Decay rate. Default is 0.9.
    - \* `eps` (float): Small constant to prevent division by zero. Default is 1e-8.
- **Method: `optimize(self, weights, biases, gradient_weights, gradient_biases)`**
  - **Description:** Updates weights and biases based on gradients.
  - **Parameters:**
    - \* `weights` (numpy array): Current weights.

- \* **biases** (numpy array): Current biases.
- \* **gradient\_weights** (numpy array): Gradients of weights.
- \* **gradient\_biases** (numpy array): Gradients of biases.
- **Returns:** Updated weights and biases.

### Class: Adam

The Adam class implements the Adam optimizer.

- **Constructor:** `__init__(self, eta=0.001, beta1=0.9, beta2=0.999, eps=1e-8)`
  - **Parameters:**
    - \* **eta** (float): Learning rate. Default is 0.001.
    - \* **beta1** (float): Exponential decay rate for the first moment estimates. Default is 0.9.
    - \* **beta2** (float): Exponential decay rate for the second moment estimates. Default is 0.999.
    - \* **eps** (float): Small

constant to prevent division by zero. Default is 1e-8.

- **Method:** `optimize(self, weights, biases, gradient_weights, gradient_biases)`
  - **Description:** Updates weights and biases based on gradients.
  - **Parameters:**
    - \* **weights** (numpy array): Current weights.
    - \* **biases** (numpy array): Current biases.
    - \* **gradient\_weights** (numpy array): Gradients of weights.
    - \* **gradient\_biases** (numpy array): Gradients of biases.
  - **Returns:** Updated weights and biases.

### Example Usage

```
import numpy as np
from Optimizers import SGD, Adadelta, Adam

# Example weights and gradients
weights = np.random.rand(3, 2)
biases = np.random.rand(2)
gradient_weights = np.random.rand(3, 2)
gradient_biases = np.random.rand(2)

# SGD Optimizer
sgd = SGD(learning_rate=0.01, momentum=0.9)
updated_weights, updated_biases = sgd.optimize(weights, biases, gradient_weights, gradient_biases)
print("SGD updated weights:", updated_weights)
print("SGD updated biases:", updated_biases)

# Adadelta Optimizer
adadelta = Adadelta(gamma=0.9)
updated_weights, updated_biases = adadelta.optimize(weights, biases, gradient_weights, gradient_biases)
print("Adadelta updated weights:", updated_weights)
print("Adadelta updated biases:", updated_biases)

# Adam Optimizer
adam = Adam(eta=0.001)
updated_weights, updated_biases = adam.optimize(weights, biases, gradient_weights, gradient_biases)
print("Adam updated weights:", updated_weights)
print("Adam updated biases:", updated_biases)
```